# Project ECHO: A Unified Cognitive Adaptation Framework for Emotion-Aware Human-Computer Interaction


## Abstract

Traditional adaptive interactive systems primarily scale execution parameters based on observable user performance metrics, often neglecting the underlying cognitive and affective states driving that performance. The Empathic Cognitive Hook Orchestrator (Project ECHO) introduces a software-only, closed-loop Human-Computer Interaction (HCI) middleware framework designed to probabilistically estimate latent cognitive, emotional, neuromotor, and trust-related states from non-intrusive behavioral interaction signals. By leveraging keystroke dynamics, Sigma-Lognormal neuromotor mouse analysis, and a Dual-Stream Transformer Emotion Recognition (DSTER) architecture, the framework extracts rich behavioral features without requiring specialized physiological hardware. These observations continuously update a Bayesian Partially Observable Markov Decision Process (POMDP) belief state, encompassing fifteen composite user states. The resulting probabilistic distribution informs an adaptation engine equipped with Human-Centered Explainable AI (HCXAI) to yield bounded, explainable interaction interventions across gaming, educational, and professional interfaces. The framework establishes a unified methodology for state-aware adaptive computing, prioritizing user autonomy and privacy through localized behavioral feature extraction.



## Keywords

Human-Computer Interaction, Affective Computing, Behavioral Biometrics, Dynamic Difficulty Adjustment, Bayesian POMDP, HCXAI, DSTER, Cognitive Load Estimation, Sigma-Lognormal Neuromotor Model.


## Executive Summary

Project ECHO addresses a critical limitation in current adaptive computing architectures: the reliance on discrete performance markers, such as failure rates or task completion times, rather than the continuous, latent cognitive and emotional realities of the user. In conventional systems, a user may fail a task due to high cognitive load, frustration, fatigue, or deliberate experimentation. Performance telemetry alone cannot disambiguate these highly divergent states, leading to inappropriate system interventions.

Conventional adaptive systems are fundamentally insufficient because they lack a probabilistic understanding of the underlying causes of user behavior. Furthermore, systems that do track cognitive states often depend on intrusive physiological hardware, such as electroencephalography (EEG) arrays or pupillometry trackers, limiting their real-world

scalability and deployment. Project ECHO solves this by utilizing universally available, non-intrusive software interaction signals. Specifically, the framework analyzes keystroke intervals and mouse movement trajectories to extract behavioral features. Within this framework, human cognition and trust are modeled as latent, partially observable states within a Bayesian POMDP, capturing the inherent epistemic and aleatoric uncertainty of human behavior.

ECHO estimates these states through an integrated mathematical pipeline comprising Cognitive Load Correlation (CLC), Sigma-Lognormal movement analysis, and a specialized DSTER Transformer model. The resulting belief state dictates adaptation policies, which include user interface simplification, dynamic difficulty adjustment (DDA), and adaptive hinting. These adaptations are transparently communicated to the user via a Human-Centered Explainable AI (HCXAI) layer.

Crucially, Project ECHO is engineered as a purely software-based HCI middleware. The current implementation successfully executes keystroke and mouse capture, feature extraction, DSTER inference, POMDP belief updating, and adaptation modulation, complete with session testing simulators. Production game engine software development kits (SDKs) and large-scale real-user dataset validations represent the primary future research trajectory, clearly distinguishing the current working prototype from proposed future expansions. The central research contribution lies in the integration of these disparate components into a singular, probabilistic, and reusable closed-loop architecture.

# Introduction

The evolution of Human-Computer Interaction has steadily shifted from rigid, user-to-interface command structures toward highly adaptive systems capable of anticipating and responding to user needs. However, realizing the true potential of cognitive-state-aware computing requires systems to infer the nuances of how a user processes information, moving beyond traditional performance-aware computing to human-state-aware computing.

Affective computing and behavioral biometrics present a robust pathway to this paradigm by analyzing subtle variations in interaction data. Historically, fields such as dynamic difficulty adjustment, human-AI trust calibration, and explainable adaptive systems have operated in isolated academic silos. Project ECHO represents a fundamental transition in interaction loops. In traditional computing, the interaction is strictly unidirectional from the user to the interface. In standard adaptive computing, the user generates behavioral data which prompts a direct system adaptation. Project ECHO reconceptualizes this loop entirely: the user generates behavioral signals, which drive latent state estimation, forming a probabilistic belief, which in turn informs an adaptive decision, culminating in an explainable interface adjustment presented back to the user. This multi-stage pipeline ensures that adaptations are empathetic, context-aware, and mathematically grounded.

# Background and Context

The demand for personalized user experiences in interactive applications, ranging from video games to complex Integrated Development Environments (IDEs), is growing rapidly. While modern game engines frequently utilize dynamic difficulty adjustment to modify enemy spawn rates or health scarcity, these modifications are typically reactionary rules based on superficial player metrics such as death counts or clear times. Meanwhile, in enterprise and

decision-support systems, a lack of trust calibration frequently leads to automation bias or algorithm aversion, scenarios where users inappropriately rely on or prematurely discard AI assistance. Project ECHO situates itself at the intersection of these domains, serving as a cognitive adaptation middleware applicable to any software environment capable of streaming standard input telemetry.

# Problem Statement

Existing adaptive interactive systems primarily depend on observable performance metrics and predefined, deterministic adaptation rules, while the user's underlying cognitive, emotional, and trust conditions remain largely unobserved. This architectural blind spot results in poorly personalized interactions, inappropriate difficulty spikes or drops, user information overload, and uncalibrated human-AI interaction dynamics.

# Motivation

The primary motivation for Project ECHO is to bridge the "emotion gap" and "cognitive gap" entirely within software-only environments. Hardware-dependent affective computing solutions, such as EEG arrays, galvanic skin response sensors, or specialized ocular trackers, face insurmountable friction regarding consumer adoption and privacy concerns. Developing a robust, software-only framework that accurately models latent states via non-intrusive biometrics is essential to democratize cognitive-aware computing across everyday devices.

# Research Objectives

The research and development of Project ECHO are guided by five primary objectives. The first objective is to formulate a unified software architecture for non-intrusive behavioral feature extraction. The second is to develop and validate mathematical models mapping interaction latency and neuromotor variables to cognitive load. The third is to implement a temporal-aware emotion inference engine, specifically the DSTER model, using keystroke dynamics. The fourth objective involves constructing a Bayesian POMDP to manage uncertainty across a multidimensional user state space. The final objective is to design a bounded adaptation engine that scales application parameters safely and explainably through HCXAI.

# Research Questions

This documentation addresses three critical research questions. First, how can non-intrusive interaction signals from standard keyboards and mice accurately proxy latent cognitive and emotional states without relying on physiological hardware sensors? Second, how does a Bayesian POMDP framework improve the stability and appropriateness of dynamic difficulty adjustment compared to deterministic, performance-based rule engines? Third, what HCXAI strategies most effectively maintain user agency and appropriate reliance when an application autonomously adapts its interface?

# Scope

The scope of Project ECHO encompasses the theoretical formulation, architectural development, simulation, and evaluation of a software-based middleware framework. It comprehensively includes the data extraction pipelines, artificial intelligence inference layers, probabilistic state trackers, and adaptation modulators. Direct, hardcoded integration with commercial AAA game engines such as Unity or Unreal via production SDKs falls outside the immediate implementation scope, but is rigorously defined within the proposed architectural roadmap.

# Assumptions

The foundational assumptions of this framework dictate that user interaction with standard peripherals contains statistically significant variance correlated with internal cognitive and affective states, as supported by behavioral biometric literature. Furthermore, the framework assumes that human cognitive load and trust can be effectively approximated as discrete latent states within a Markovian process. Finally, it assumes that the application domains connecting to the ECHO middleware possess tunable parameters suitable for continuous, programmatic modulation.

# System Positioning

Project ECHO is explicitly positioned as an HCI middleware framework and a behavioral-state inference engine. It functions as a probabilistic user-state estimation system and a cognitive adaptation engine designed for closed-loop human-centered adaptive computing.
Conversely, Project ECHO is strictly defined by what it is not. It is not a hardware project, a robotics project, or a physical actuator controller. It is not merely a standalone game, a generic chatbot, or a static DDA script. Most importantly, it is not a system intended to remove human decision authority. The guiding philosophy of the architecture is that Project ECHO adapts the interface and the interaction experience, not the user's fundamental autonomy.

# Literature Review

To contextualize ECHO, current research is categorized into major foundational pillars, moving away from fragmented analysis toward a cohesive understanding of affective computing and adaptive systems.

## A. Adaptive Game Systems & B. Dynamic Difficulty Adjustment

Traditional dynamic difficulty systems utilize reinforcement learning techniques such as Proximal Policy Optimization (PPO), Deep Q-Networks (DQN), or Fuzzy Logic to balance gameplay. These systems typically track parameters like damage dealt, kill speed, and death counts. While these methods excel at providing a balanced mechanical challenge, they are fundamentally limited by their reliance on observable performance, ignoring the emotional reality of the player. Project ECHO adopts the concept of continuous parameter modulation from this domain but entirely shifts the triggering mechanism from mechanical performance metrics to probabilistically estimated cognitive and affective states.

## C. Player Modeling & D. Cognitive Workload Estimation

Existing cognitive workload systems heavily rely on physiological metrics. Research demonstrates high classification accuracy using pupil diameter, electrodermal activity, and electromyography. However, these approaches share a critical limitation: the absolute necessity of specialized biometric hardware. Project ECHO recognizes the validity of tracking cognitive load but substitutes physiological hardware with behavioral proxies, treating cognition as a latent variable inferred solely through software interaction patterns.

## E. Behavioral Biometrics, F. Keystroke Dynamics, & G. Mouse Movement Modeling

The Kinematic Theory of rapid human movements and its associated Sigma-Lognormal model have been effectively applied to decompose complex writing and pointing movements into primitives characterized by lognormal velocity profiles. This model accurately parameterizes the neuromuscular impulse response. ECHO adopts the Sigma-Lognormal model to extract neuromotor strain and signal-to-noise ratio from standard mouse trajectories, mapping physical execution degradation to cognitive state without requiring a digital stylus or specialized tablet.

## H. Emotion Recognition, I. Multimodal Emotion Recognition, & J. Transformer-Based Affective Computing

Recent advances in emotion recognition utilize Dual-Stream Transformers to classify emotion from keystroke dynamics, effectively extracting temporal and contextual data to yield high F1 scores. A common limitation in multimodal architectures is modality collapse or unmanageable latency during real-time inference. Project ECHO natively integrates the DSTER paradigm to ensure temporal affect tracking remains lightweight and highly accurate based solely on typing cadence.

## K. Bayesian Inference & L. POMDP-Based Human Modeling

Bayesian inference, including Kalman filters and stochastic architectures, excels at modeling epistemic and aleatoric uncertainty in noisy data environments. Partially Observable Markov Decision Processes are ideal for scenarios where the true state of the environment—or in this case, the human mind—cannot be directly measured. ECHO leverages a POMDP to track the hidden user states of Trust and Cognitive Load based on noisy biometric observations, ensuring that system adaptations do not wildly oscillate based on a single erratic keystroke.

## M. Human-AI Trust, N. Trust Calibration, O. Explainable AI, & P. Human-Centered Explainability

Trust calibration aims to mitigate algorithm aversion and automation bias, ensuring the user relies on the system appropriately. Human-Centered Explainable AI transitions traditional XAI, which asks why the model made a decision, to human-centric presentation, asking what the user should know based on their current cognitive load. ECHO incorporates Trust directly into the POMDP state matrix and utilizes HCXAI to modulate explanation verbosity dynamically.

### Q. Adaptive User Interfaces, R. HCI Personalization, S. Edge/Real-Time AI, & T. Privacy-Preserving Analytics

Edge deployment limits inference latency, while privacy demands strictly localized feature extraction. Project ECHO addresses these requirements by preprocessing all behavioral data locally, immediately discarding raw event logs to ensure user privacy, and operating within strict real-time compute budgets.

# Literature Synthesis

The analysis of the literature demonstrates profound, yet deeply isolated, advancements. High-accuracy emotion recognition via keystrokes using advanced transformer architectures exists entirely separately from Bayesian POMDP decision-making utilized in robotics. Furthermore, both are isolated from the Sigma-Lognormal neuromotor analysis utilized in kinesiology, which is separate from the dynamic difficulty adjustment paradigms found in commercial gaming. The literature lacks a unified, software-only framework that synthesizes these independent breakthroughs into a singular operational pipeline.

# Research Gap

Existing literature consistently studies emotion recognition, cognitive workload estimation, player modeling, trust estimation, dynamic difficulty adjustment, and explainability as fragmented, discrete challenges. Consequently, systems are either highly accurate but require laboratory-grade hardware, or they are software-only but rely on shallow, heuristic-based performance tracking. No previous work has integrated a software-only behavioral biometric pipeline—combining DSTER and Sigma-Lognormal analysis—with a Bayesian POMDP to drive bounded, HCXAI-supported adaptive interfaces.

Project ECHO's central contribution is therefore defined by integration. The framework unites probabilistic state estimation, closed-loop adaptation, software-only interaction sensing, human-centered explainability, bounded adaptation, and a reusable middleware architecture into a single coherent system.

# Proposed Solution

Project ECHO proposes a unified middleware that bridges the substantial gap between raw human interaction and state-aware application adaptation. The framework captures raw peripheral data, translates it into statistical and neuromotor features, executes deep learning inference to assess affect, probabilistically updates a multi-dimensional belief state, and routes this state to a rule-based adaptation engine. This engine subsequently feeds dynamic, explainable updates back to the connected application, creating a continuous, empathetic feedback loop.

# Overall Architecture

The logical architecture of Project ECHO represents a unified closed-loop framework ensuring uninterrupted data and control flow. The architecture strictly adheres to the following sequence:

USER \rightarrow NON-INTRUSIVE SOFTWARE INTERACTION SIGNALS \rightarrow INPUT CAPTURE \rightarrow PREPROCESSING \rightarrow SIGNAL QUALITY VALIDATION \rightarrow BEHAVIORAL FEATURE EXTRACTION \rightarrow COGNITIVE LOAD ESTIMATION \rightarrow NEUROMOTOR / MOUSE MOVEMENT ANALYSIS \rightarrow DSTER EMOTION INFERENCE \rightarrow MULTI-SIGNAL FUSION \rightarrow BAYESIAN STATE ESTIMATION \rightarrow POMDP BELIEF STATE \rightarrow TRUST + COGNITIVE STATE ESTIMATION \rightarrow ADAPTATION POLICY \rightarrow DYNAMIC UI / ASSISTANCE / DDA / EXPLANATION \rightarrow HCXAI \rightarrow ADAPTIVE INTERFACE / GAME / APPLICATION \rightarrow NEW USER INTERACTION \rightarrow CONTINUOUS FEEDBACK LOOP.

# Detailed System Workflow

Data flows continuously and asynchronously through the architecture. The input capture modules gather raw OS-level events. Preprocessing sanitizes timestamps and filters out hardware interrupts. Feature engineering splits the data conceptually: typing data feeds the cognitive load correlation estimators and digraph latency modules, while spatial mouse data feeds the Sigma-Lognormal neuromotor analysis. The signal-to-noise ratio validator subsequently gates poor-quality signals to prevent model poisoning. Quality-assured features then pass to the DSTER module for affect inference. The multidimensional output, comprising cognitive load metrics, signal quality, typing latency, emotion labels, and inference confidence, acts as the observation vector in the Bayesian POMDP module. The resulting belief state defines the current probability distribution of the user's cognitive load and trust levels. This triggers the adaptation engine and parameter modulator to calculate safe, bounded application changes, which are finally contextualized by the HCXAI formatter before altering the user experience.

# Mathematical Foundations

This section rigorously defines the mathematical models propelling the ECHO architecture, mapping theoretical constructs to implementation realities.

## 10.1 Inter-Key Interval

The Inter-Key Interval (IKI) represents the temporal latency between consecutive keystrokes, providing a foundational behavioral feature for cognitive tracking. The mathematical expression is defined as:
IKI_i = t_{i+1} - t_i
The variable t_i represents the exact microsecond timestamp of the i-th key press event, while t_{i+1} represents the timestamp of the subsequent key press. The intuition behind this metric is that typing rhythm is highly susceptible to cognitive interference. By aggregating IKI over temporal sliding windows, ECHO extracts the mean IKI, median IKI, overall variance, standard deviation, and the coefficient of variation. Increased variance and extended pause durations serve as primary indicators of hesitation or high cognitive demand. The implementation calculates these metrics in real-time, feeding the data forward to the feature extraction layer. A limitation of this model is that hardware polling rates and operating system-level interrupts can

introduce micro-latency noise, necessitating robust preprocessing.

## 10.2 Digraph Latency

Digraph latency extends the base IKI concept by analyzing the timing between specific alphanumeric key pairs to capture temporal behavior within and between linguistic components. Consecutive-key timing captures temporal behavior by distinguishing between automatic neuromuscular execution and active cognitive formulation. Within-word timing typically relies on ingrained muscle memory, indicating low cognitive demand. Conversely, between-word timing correlates with lexical retrieval, sentence formulation, and active decision-making, signifying higher cognitive demand. By tracking pauses, hesitations, and temporal variability across specific digraphs, ECHO gathers behavioral evidence of cognitive load. It is critical to note that ECHO does not overclaim that digraph latency directly diagnoses complex human emotion; rather, it is strictly treated as behavioral evidence that feeds into broader inferential models.

## 10.3 Cognitive Load Correlation

Project ECHO measures the relationship between the complexity of the content being interacted with and the physical latency of the interaction itself to estimate cognitive burden. The relationship is defined as:
CLC = \rho_s(R_C, R_{IKI})
The variable \rho_s represents the Spearman rank correlation coefficient. The variable R_C denotes the ranked content-complexity measure, such as the lexical density or interface complexity of the task, while R_{IKI} denotes the ranked inter-key latency. The Spearman correlation is utilized because it effectively captures monotonic, non-linear relationships, making the metric exceptionally robust against outlier pauses that would otherwise skew Pearson correlations. ECHO's nominal configurations establish specific interpretation thresholds. A CLC value less than 0.35 suggests potential mechanical, copy-paste, or low-generative behavior. A value between 0.35 and 0.55 indicates normal active composition. A CLC greater than 0.55 flags increased deliberation, unfamiliarity, or high cognitive demand. These thresholds represent design configuration values and require independent experimental validation for different demographic cohorts.

## 10.4 Spearman Correlation

The mathematical formulation for the Spearman rank correlation used in the CLC estimation is:
\rho_s = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}
The variable d_i represents the numerical difference between the two ranks (R_{C,i} - R_{IKI,i}) for each observed interaction event. The variable n represents the total number of observations within the temporal window. In the event of tied ranks, standard averaging procedures are applied. The output is strictly bounded between -1 and 1, providing a normalized metric that feeds directly into the state fusion module, limiting unbounded variance in downstream calculations.

## 10.5 Sigma-Lognormal Neuromotor Model

The Sigma-Lognormal model decomposes complex, rapid human movements—specifically mouse trajectories within ECHO—into a vectorial summation of primitive strokes governed by

lognormal velocity profiles. The general model for a lognormal velocity profile is expressed as:
$$|\vec{v}(t)| = \sum_{i=1}^{n} \frac{D_i}{\sqrt{2\pi}\sigma_i(t-t_{0i})} \exp\left(-\frac{(\ln(t-t_{0i}) - \mu_i)^2}{2\sigma_i^2}\right)$$
The variable $D_i$ represents the amplitude of the i-th motor primitive. The variable $t_{0i}$ dictates the time of occurrence or onset of the primitive. The variable $\mu_i$ represents the log-time mean, and $\sigma_i$ is the log-time standard deviation. The intuition is that complex movement is not generated as a single continuous signal but rather as a sum of discrete, overlapping motor commands. ECHO implements a parameter extraction and optimization procedure to fit these theoretical primitives to the raw mouse coordinates, generating a reconstructed velocity profile. High variance in $\sigma_i$ or poorly synchronized onset sequences ($t_{0i}$) serve as powerful indicators of neuromotor stress, fatigue, or cognitive interference.

## 10.6 Signal-to-Noise Ratio

To ensure behavioral data is actionable, a Signal-to-Noise Ratio (SNR) calculation assesses the reconstruction quality of the Sigma-Lognormal model against the raw input trajectory. Signal is defined as the mathematical reconstruction of the velocity profile, while noise is defined as the residual variance between the raw input and the reconstruction. Represented logarithmically in decibels (dB), this metric acts as an essential quality gate. ECHO utilizes strict project thresholds: an SNR below 15 dB is classified as poor, resulting in the data being gated or assigned extremely low confidence. An SNR between 20 and 25 dB is classified as good, while anything exceeding 25 dB is considered excellent. These act as configuration thresholds rather than universally valid scientific constants, ensuring that subsequent Bayesian updates are not corrupted by noisy, unusable data.

## 10.7 DSTER

The Dual-Stream Transformer Emotion Recognition (DSTER) architecture mathematically extracts affect from behavioral sequence embeddings. The architecture relies on two parallel processing mechanisms: Temporal-over-Channel processing, which emphasizes the rhythmic and sequential patterns of interaction, and Channel-over-Temporal processing, which models correlations among the distinct features within the input vector. Embeddings are generated from the keystroke sequences and infused with positional representations to retain temporal order. Both streams utilize standard Transformer self-attention formulations:
$$\text{Attention}(Q,K,V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$
The matrices Q (Query), K (Key), and V (Value) are derived from the input embeddings, while the scaling factor $\sqrt{d_k}$ prevents vanishing gradients during softmax computation. The temporal and channel aggregations are eventually subjected to stream fusion before passing through a dense classifier.

## 10.8 Softmax Emotion Probability

Following DSTER stream fusion, the final classification layer outputs a probability distribution over the predefined emotion labels.
$$P(y=k|x) = \frac{e^{z_k}}{\sum_j e^{z_j}}$$
The variable $z_k$ represents the raw logit output for emotion class k. The implementation supports five distinct labels: frustrated, focused, fatigued, neutral, and engaged. Crucially, the model outputs not just the emotion label and its probability, but a strict confidence score. ECHO

relies heavily on uncertainty; if the confidence score is below 0.6, the system will actively suppress aggressive adaptations, recognizing the inherent ambiguity of emotion recognition.

## 10.9 Bayesian State Estimation

Because human cognitive states cannot be directly observed, ECHO employs a complete Bayesian formulation to maintain a probabilistic belief state over time. The update formulation is defined as:
$b'(s') = \eta O(o|s') \sum_{s} T(s'|s)b(s)$
The variable $b(s)$ represents the prior belief distribution over the state space. $T(s'|s)$ is the transition model, defining the probability of moving from state $s$ to state $s'$. $O(o|s')$ represents the likelihood of receiving observation $o$ given the user is in state $s'$. The variable $\eta$ is a necessary normalization constant ensuring the updated probabilities sum to one. The process flows sequentially from prediction to observation ingestion, likelihood calculation, state updating, and final normalization.

## 10.10 POMDP

Project ECHO formally maps the user interaction environment as a Partially Observable Markov Decision Process. The POMDP is defined by the tuple:
$P = (S,A,O,T,Z,R,\gamma)$
The set $S$ defines the latent user states. $A$ represents the set of adaptation actions available to the system. $O$ constitutes the behavioral observations. $T$ is the transition model. $Z$ represents the observation model (equivalent to $O(o|s')$). $R$ is the reward function guiding optimal policy, and $\gamma$ is the discount factor prioritizing immediate versus long-term user stability. The state space ($S$) is explicitly defined by two parallel variables. Trust comprises five levels: Very Low, Low, Medium, High, and Very High. Cognitive Load comprises three levels: Low, Medium, and High. Consequently, ECHO operates across a composite state space where $|S| = 5 \times 3 = 15$ explicit states.

## 10.11 Belief State

The belief state $b_t(s)$ functions as a continuous probability distribution over the 15 hidden user states. This probabilistic approach is vastly superior to deterministic classification because human behavior is inherently noisy. For example, instead of a rigid rule declaring the user is "Frustrated," the belief state reflects a nuanced reality: $P(\text{Frustrated}) = 0.62$, $P(\text{Focused}) = 0.18$, $P(\text{Fatigued}) = 0.12$. A deterministic system might trigger an immediate, massive difficulty reduction based on the primary label, whereas the belief state tempers the response by acknowledging a 38% probability that the user is not actually frustrated, leading to a more measured, bounded adaptation.

## 10.12 Trust Modeling

Within ECHO, trust is explicitly modeled as a dynamic state rather than a static user trait. It directly addresses the HCI challenges of automation bias (excessive trust) and algorithm aversion (insufficient trust), aiming instead for appropriate reliance. Trust transitions dynamically over time based on system reliability and task risk. The current implemented representation manages trust implicitly through the POMDP transition matrices. A theoretically appropriate

future formulation involves explicit mathematical updates, penalizing the trust state when a user actively rejects a system recommendation or manually overrides an adaptation, thereby explicitly linking behavioral pushback to the probabilistic belief of trust.

## 10.13 State Fusion

The framework must reconcile disparate data streams. The state fusion module aggregates the CLC, IKI metrics, Sigma-Lognormal SNR, DSTER emotion vectors, emotion confidence, and neuromotor indicators to generate the discrete observation vector o. ECHO currently implements confidence-weighted late fusion. If the SNR is poor or the DSTER confidence is low, the observation likelihood O(o|s') is flattened, heavily weighting the prior belief b(s) and preventing erratic state shifts. Proposed future work includes transitioning to attention-based Bayesian fusion for more granular, continuous feature weighting.

## 10.14 Adaptation Function

The adaptation engine translates the belief state into actionable HCI interventions. The function is formally described as:
$A_t = \pi(b_t, r_t, c_t)$
The variable $A_t$ is the output action. The function $\pi$ represents the adaptation policy. The variable $b_t$ is the current belief state. The variable $r_t$ represents the task risk or context severity, and $c_t$ denotes the specific constraints of the application or game context. Adaptation output categories are broad, encompassing response verbosity adjustments, UI simplification, uncertainty presentation, graduated hinting, difficulty modulation, target highlighting, HUD decluttering, and opt-in assistance prompts.

## 10.15 Parameter Modulation

To prevent adaptation instability, raw parameter recommendations must be modulated. The flow moves from raw adaptation to sensitivity scaling, followed by mathematical smoothing, the application of hard bounds, and finally, the output of the final parameter. Exponential smoothing is utilized to govern state transitions:
$P_{new} = \alpha P_{target} + (1-\alpha)P_{old}$
The variable $\alpha$ acts as a sensitivity scaling factor. Proposed future formulations include adaptive rate limiting based on the derivative of the belief state to further suppress oscillatory "yo-yo" difficulty adjustments.

## 10.16 Hint Engine

ECHO implements a graduated hint engine to preserve user autonomy. The current implemented timing rationale utilizes strict temporal graduation: 10 seconds triggers a subtle nudge, 20 seconds triggers a tactical guide, and 30 seconds reveals the explicit solution. This ensures that assistance is progressive rather than immediate, aligning with pedagogical best practices. Future improvements propose dynamically mapping the cognitive state and confidence metrics to the hint timer. A high-probability belief of cognitive overload would mathematically reduce the temporal thresholds, delivering assistance more rapidly when struggle is highly certain.

# Software Architecture & Module-Level Implementation

The actual project implementation is heavily documented at the module level, confirming its software-only architecture.

| Module Directory | File Name | Purpose and Processing Overview |
|---|---|---|
| input_capture/ | keystroke_listener.py | Interfaces with OS-level hooks to silently capture interaction timing without parsing semantic text. Computes base latencies. |
| input_capture/ | mouse_listener.py | Tracks continuous X/Y coordinate trajectories and raw velocity. |
| feature_extraction/ | digraph_latency.py | Calculates temporal variances between specific key pairings. |
| feature_extraction/ | clc_estimator.py | Implements the Spearman rank correlation to map typing speed against application content complexity. |
| feature_extraction/ | sigma_lognormal.py | Extracts neuromotor primitives and parameters (D, t_0, \mu, \sigma) from raw mouse coordinates. |
| feature_extraction/ | snr_validator.py | Computes the signal-to-noise ratio in decibels to validate lognormal reconstructions. |
| game_framework/ | session_simulator.py | Generates synthetic behavioral data streams to test the POMDP without requiring a live game engine. |
| deployment/ | quantization.py | Contains utilities for INT8 and FP16 model quantization to ensure rapid edge deployment. |

These modules feature highly decoupled dependencies, ensuring that a failure in the sigma_lognormal.py extraction merely lowers observation confidence rather than crashing the primary adaptation_engine loop.

# Implemented vs. Proposed Status

It is critical to demarcate the current engineering reality of Project ECHO from its theoretical roadmap.

| Component Category | Implementation Status |
|---|---|
| Keystroke & Mouse Capture | **Implemented** |
| IKI Extraction & CLC Estimation | **Implemented** |
| Sigma-Lognormal Processing | **Implemented** |
| SNR Validation | **Implemented** |

| Component Category | Implementation Status |
|---|---|
| DSTER Architecture & Emotion Classification Interface | **Implemented** |
| Bayesian POMDP & Belief Updating | **Implemented** |
| Trust Representation & Cognitive Load Fusion | **Implemented** |
| Adaptation Engine & Hint Engine | **Implemented** |
| DDA Parameter Modulation & Session Simulation | **Implemented** |
| Visualization & Automated Tests | **Implemented** |
| Real-User Dataset Validation | *Not Demonstrated / Future* |
| Production Unity & Unreal SDKs | *Not Demonstrated / Future* |
| Real Commercial-Game Integration | *Not Demonstrated / Future* |
| Large Human-Subject Evaluation | *Not Demonstrated / Future* |
| Production-Scale Deployment | *Not Demonstrated / Future* |

# Software-Only Requirement

The core Project ECHO implementation is unequivocally software-only. The core input mechanisms rely exclusively on standard keyboard interaction, mouse movement trajectories, software-generated telemetry, and derived mathematical behavioral features. There is absolute zero dependency on Arduino microcontrollers, ESP32 chips, Jetson Nano devices, Raspberry Pi hardware, servo motors, robotics, physical actuators, or EEG hardware. While optional computer-vision modalities may be discussed as theoretical extensions, they do not form the methodological backbone of the current implementation.

# Dynamic Difficulty Adjustment (DDA)

ECHO functions as a powerful DDA middleware by simulating adjustments across specific game profiles.

- **Left 4 Dead:** Adaptation parameters target the Director AI, modulating enemy spawn rates, horde frequency, health-item scarcity, special-infected caps, and wanderer density. A high cognitive load state reduces horde frequency to alleviate cognitive strain.
- **Alien: Isolation:** Parameters include menace level, tension budget, search radius, vent probability, and noise sensitivity. High user frustration limits the tension budget, granting the player brief psychological reprieves.
- **Resident Evil 4:** Modulates enemy accuracy, enemy aggression, item drop quality, damage multipliers, and the core action score.
- **Silent Hill: Shattered Memories:** Adjusts world decay rates, monster bias, puzzle obfuscation, lighting, and static interference to maintain atmospheric tension without inducing absolute despair. Safe parameter ranges and exponential smoothing algorithms dictate the adaptation frequency to explicitly mitigate the risk of jarring difficulty oscillation. Currently, ECHO presents existing functionality as parameter simulation and profile support; it does not directly modify commercial retail game executables without future SDK implementation.

# Session Simulation

The session_simulator.py is utilized to robustly test the state changes, belief updates, emotion confidence handling, and parameter changes without requiring active human playtesting.

- **Natural Scenario:** Simulates standard fluctuations in performance, validating that ECHO's smoothing algorithms ignore minor noise.
- **Struggle Scenario:** Injects high latency, low SNR, and DSTER frustration markers. Validates that the POMDP swiftly transitions to a high cognitive load state, triggering substantial difficulty reduction.
- **Mastery Scenario:** Simulates rapid, precise typing and high DSTER engagement. Validates upward parameter scaling.
- **Volatile Scenario:** Injects wildly erratic user states to verify that ECHO maintains stability and bounded adjustments, preventing chaotic gameplay loops.

# Testing Methodology

Project ECHO requires an exhaustive testing strategy.

- **Unit Testing:** Validates independent mathematical models, passing controlled synthetic arrays into the Spearman correlation and Sigma-Lognormal modules to verify expected deterministic outputs.
- **Integration Testing:** Ensures the sequential pipeline successfully hands data from input_capture through feature_extraction, into the DSTER and POMDP modules, resulting in a synthesized adaptation parameter.
- **Simulation Testing:** Leverages the session_simulator.py to evaluate the system over thousands of synthetic epochs.
- **Performance Testing:** Measures the inference latency, ensuring the end-to-end cycle processes within strict sub-100 millisecond thresholds to remain imperceptible to the user.
- **Robustness & HCI Testing:** Future metrics require evaluating the stability of the parameter outputs under extreme noise and conducting comprehensive user study validations.

# Current Testing Status

Execution of the core repository test suite yields verifiable metrics: **59 tests passed, 3 failed, and 11 resulted in errors**. The primary cause of the inconsistency is a mapping drift within the integration modules. The automated test files expect generic category designations (e.g., fps, puzzle, strategy, racing, survival_horror). However, the operational game_registry.py utilizes specific title profiles (left_4_dead, alien_isolation, resident_evil_4, silent_hill). This mismatch causes integration tests to fail when fetching nonexistent generic profile keys. The proposed fix involves standardizing the canonical game registry to support mapped aliases, ensuring robust regression testing in the future without hiding the current integration faults.

# Machine Learning Evaluation

To validate the behavioral biometrics and DSTER models, a complete machine learning evaluation framework must encompass Accuracy, Precision, Recall, Macro F1, Weighted F1, ROC-AUC, Confusion Matrices, and Calibration error metrics. Given the sequential nature of behavioral interactions, the models must be tested for temporal generalization (performance stability over long sessions) and cross-session generalization. Crucially, future real-data studies

must employ Leave-One-User-Out Cross-Validation (LOUO-CV) to prove that the emotion classifiers are learning universal behavioral traits rather than memorizing the idiosyncratic typing quirks of a specific training user.

# Behavioral Model Evaluation

The continuous behavioral signals are evaluated using appropriate statistical distributions. Evaluations encompass IKI distribution analyses, digraph latency clustering, CLC distribution plotting, Sigma-Lognormal reconstruction error rates, and movement variability metrics. Appropriate statistical tests include the Spearman correlation for CLC. For future comparative analyses across static versus adaptive interface conditions, repeated-measures ANOVA (where parametric assumptions hold) or the Kruskal-Wallis test (for non-parametric data) will be utilized to rigorously determine effect sizes and confidence intervals.

# Adaptation Evaluation

Adaptive interfaces introduce significant instability risks. ECHO evaluates DDA via explicit mathematical metrics:

- **Adaptation Reversal Rate (ARR):** Defined as the number of directional reversals divided by the total adaptation events. A high ARR indicates frustrating oscillation.
- **Parameter Overshoot (O):** Defined as the maximum excursion beyond the desired parameter target.
- **Convergence Latency:** Measures the number of epochs required for the system to stabilize parameters following a major shift in the user's belief state.

# Human-Computer Interaction Evaluation

A rigorous HCI evaluation is proposed to validate the efficacy of ECHO in real-world scenarios. The experimental design relies on three distinct conditions:

- **Condition A:** Static interface (no adaptation).
- **Condition B:** Traditional performance-based DDA (adapting solely to mechanical game performance).
- **Condition C:** Project ECHO (cognitive and affective state-based adaptation). Metrics measured will include task completion time, error rates, cognitive workload assessed via the NASA-TLX index, System Usability Scale (SUS) scores, user-reported trust, and retention rates. The experiment will require careful sample size considerations, strict randomization, and counterbalancing to mitigate learning effects and fatigue.

# Real User Dataset Proposal

To advance the research beyond synthetic validation, a comprehensive future dataset must be established. The proposed schema encompasses critical fields: session_id, user_id, task_difficulty, timestamp, key_event, IKI, digraph_latency, mouse_trajectory, CLC, SNR, emotion_label, confidence, cognitive_load_rating, trust_rating, hint_accepted, manual_override, and completion_time. Collection must prioritize explicit consent, strict anonymization of typing semantics, and rigorous quality control regarding the test/train splits.

# Personalization

Universal thresholds—such as assuming a CLC > 0.55 universally indicates high cognitive load—are fundamentally insufficient due to massive variances in baseline human performance. A user calibration process is required: an initial baseline session must record natural interaction, allowing for behavioral normalization (e.g., z-score normalization) to establish personalized thresholds. These thresholds feed into the adaptive inference models, enabling personalized Bayesian priors and continual online learning to track user skill mastery over time. These represent mathematically justified future approaches distinct from the currently implemented universal thresholds.

# Fairness and Robustness (Context: Chennai, Tamil Nadu, India)

Behavioral biometrics carry inherent risks regarding demographic bias. When contextualizing Project ECHO for deployment in Chennai, Tamil Nadu, India, fairness testing is paramount.

- **Linguistic Variance:** Users frequently switch between English, formal Tamil, and Tanglish. Typing speeds and digraph latencies will shift dramatically based on language familiarity and localized keyboard layouts (e.g., Tamil 99 versus standard QWERTY). The model must not penalize users for phonetic transliteration pauses, conflating linguistic switching with cognitive overload.
- **Hardware and Infrastructure:** Users utilizing low-tier hardware or inconsistent network infrastructures may exhibit irregular mouse polling rates, resulting in low Sigma-Lognormal SNR. The system must degrade gracefully, relying heavily on the Bayesian prior belief state rather than issuing erratic, uncalibrated adaptations based on false negatives.

# Privacy and Security

Keyboard and mouse telemetry constitute highly sensitive behavioral biometric information. The future architecture emphasizes privacy through data minimization and local processing.

- **Data Flow:** Raw events are ingested, local feature extraction algorithms generate mathematical abstractions (e.g., IKI, SNR), and the raw event log is instantly discarded. Only the abstracted features are stored or transmitted.
- **Threat Models:** This architecture specifically mitigates behavioral profiling, data leakage, and inference attacks designed to reverse-engineer typed sequences, ensuring that no semantic information ever leaves the client machine.

# HCXAI and Explainability

Traditional Explainable AI (XAI) primarily answers the developer's question: "What features influenced the prediction?" Human-Centered Explainable AI (HCXAI) pivots to address the user: "What should the user be told, when, and at what level of detail?". Adaptive explanation policies scale based on the POMDP belief states. If the system registers a high cognitive load, presenting a dense, multi-page justification for a parameter change is counterproductive and exacerbates the load. Instead, ECHO provides minimal visual cues. Conversely, during high-risk

tasks with high trust levels, ECHO explicitly surfaces model uncertainty to preserve appropriate reliance, avoiding dangerous automation bias.

# Safety and Human Agency

Project ECHO adheres to strict safety principles governing affective computing. The framework executes no autonomous critical decisions, forces no unrequested interventions, and strictly avoids manipulative trust shaping or dark patterns. It prohibits the unnecessary exposure of sensitive biometric data. The fundamental principle is to adapt the interface, never the user's autonomy. Every adaptation is bounded, and explicit user override mechanisms must be present in any commercial deployment to prevent the direct replacement of human authority.

# Performance and Real-Time Considerations

Real-time adaptive systems must respect strict computational latency budgets. Inference latency must be minimized (targeting sub-50 ms processing times) to ensure adaptations remain imperceptible. Project ECHO manages computational complexity through batching, streaming inference across localized temporal windows, and optimizing CPU utilization. For future deployments, model quantization techniques (INT8 and FP16) and ONNX export protocols are specified to aggressively shrink the memory footprint and maximize throughput on standard consumer-grade hardware.

# Deployment

While the current repository infrastructure consists of localized Python modules and analytical scripts, future production deployment necessitates modular encapsulation. The deployment architecture delineates the use of model exporters, dynamic quantization, and edge runtimes to package the inference and POMDP engines into a lightweight, standalone background service, cleanly separated from the immediate prototype scripts.

# API / Middleware Design

To function effectively as a reusable cognitive intelligence layer, ECHO targets a proposed REST and WebSocket middleware API schema. Proposed endpoints include:

- POST /session/start: Initializes session tracking and POMDP state vectors.
- POST /events/keyboard & POST /events/mouse: Ingests raw telemetry into the local preprocessor.
- POST /infer: Triggers the DSTER and feature extraction layers.
- GET /belief: Returns the 15-state probability distribution.
- POST /adapt: Queries the adaptation policy engine for bounded parameter adjustments.
- GET /explanation: Retrieves context-aware HCXAI formatting. These represent Proposed Future APIs designed to facilitate eventual integration, rather than currently implemented network protocols.

# Game Engine Integration

Commercial integration with engines such as Unity, Unreal Engine, or Godot is proposed via a specialized SDK architecture. The data flow dictates that the Application Engine interfaces with the ECHO SDK plugin, streaming telemetry to the local feature extraction layer. The state estimation and adaptation policy modules formulate adjustments, triggering an application parameter update. This parallel architecture ensures ECHO operates as a decoupled cognitive thread, never blocking the primary rendering loops of the host application.

# Real-World Applications (Context: Chennai, Tamil Nadu, India)

ECHO's middleware extends significantly beyond video games, offering transformative applications across multiple domains.

- **Educational Platforms:** With the Tamil Nadu state government aggressively pushing digital literacy initiatives, an ECHO-integrated e-learning platform can silently track student cognitive load via typing latency during programming assignments. High sustained load triggers automated hinting, dynamically slowing the curriculum progression to match the student's Zone of Proximal Development.
- **Enterprise IDEs:** Within Chennai's major IT corridors (e.g., SIPCOT, Tidel Park), software developers face chronic cognitive fatigue. An ECHO-aware IDE can observe erratic keystrokes and low-SNR mouse movements, responding by temporarily decluttering the UI and suppressing non-critical compilation notifications, thus preserving developer focus.
- **Accessibility Systems:** For users experiencing temporary motor instability or fatigue, a drop in Sigma-Lognormal SNR triggers an adaptive increase in UI target sizes and spacing, facilitating smoother interaction without requiring explicit manual configuration.

# Limitations

As an academic research prototype, Project ECHO possesses several limitations requiring critical, honest acknowledgment. Currently, the emotion recognition and POMDP belief updates rely heavily on synthetic validation and simulated interaction logs rather than expansive, real-world human telemetry datasets. Assigning objective ground-truth labels for subjective states like "Trust" based purely on non-intrusive latency data remains challenging without robust self-reporting studies. Technically, a significant repository inconsistency exists regarding the game registry and test mismatches. Furthermore, the system lacks production Unity/Unreal SDKs, meaning it currently relies on simulated parameter output rather than genuine real-game integration. Universal threshold limitations and the lack of personalized calibration highlight the need for extensive demographic fairness testing before wide-scale deployment.

# Future Work

ECHO's evolutionary roadmap spans ten distinct, actionable phases:

1. **Repository Stabilization:** Resolve the Game Registry test mismatches and align configuration files.
2. **Real Behavioral Dataset:** Deploy an ethical telemetry capture tool to build diverse interaction baselines.
3. **Personalization:** Implement the user-calibration and robust scaling pipelines.

4. **Advanced Multimodal Fusion:** Explore integrating computer-vision cues purely as optional parallel modalities.
5. **Real Game/Application SDK:** Build native plugins for major engines and IDEs.
6. **Human-Subject Study:** Conduct the extensive Condition A/B/C empirical validation.
7. **Privacy-Preserving Local Inference:** Finalize optimized edge models for secure, offline execution.
8. **Production Deployment:** Containerize the middleware for broad distribution.
9. **Cross-Domain Validation:** Test the framework across varied domains, contrasting educational platforms against high-reliability cybersecurity dashboards.
10. **Longitudinal Adaptive User Modeling:** Track belief state shifts over multi-month usage spans to measure chronic fatigue and long-term skill mastery.

# Advanced Future Mathematical Extensions

While currently relying on discrete POMDPs and Spearman correlations, future implementations propose incorporating advanced mathematical models.

- **Bayesian Neural Networks (BNNs):** Proposed to replace standard DSTER classifiers, outputting inherent epistemic uncertainty alongside the emotion prediction, thereby seamlessly feeding the POMDP observation likelihoods.
- **Partially Observable Reinforcement Learning (PORL):** Transitioning the discrete 15-state POMDP into a continuous latent space representation, optimizing the adaptation policy via deep reinforcement learning with an explicit HCXAI constraint penalty.
- **Risk-Sensitive POMDPs:** Formally embedding Conditional Value at Risk (CVaR) into the adaptation engine to mathematically guarantee that parameter modulation never breaches a catastrophic failure state for the user.

# Formal System Model

Project ECHO operates as a unified, closed-loop dynamical mathematical system. The overarching formulation is:

1. **User behavioral input:** x_t (representing raw keystrokes and mouse coordinates).
2. **Feature extraction:** f_t = \phi(x_{t-k:t}) (yielding IKI, CLC, and Sigma-Lognormal vectors).
3. **Signal quality:** q_t = Q(f_t) (validating SNR).
4. **Emotion model:** e_t = \text{DSTER}(f_t) (incorporating signal quality gating).
5. **Belief state:** b_t = P(s_t | o_{1:t}) where the observation vector incorporates f_t, q_t, e_t.
6. **Adaptation:** a_t = \pi(b_t, r_t, c_t) (determining the bounded parameter shift).
7. **Environment state:** y_{t+1} = F(y_t, a_t, x_t) (the application updates difficulty or UI elements).
8. **Feedback:** x_{t+1} = G(y_{t+1}, u_{t+1}) (the user perceives the adaptation, generating the next behavioral input).

This formal, continuous loop—Human \rightarrow Observation \rightarrow Inference \rightarrow State \rightarrow Action \rightarrow Environment \rightarrow Human—constitutes the mathematical backbone of the entire report.

# Algorithmic Pseudocode

The following pseudocode algorithms systematically map the implementation logic of Project

ECHO.

**Algorithm 1: Behavioral Data Acquisition** Initialize system-level hooks for Keyboard and Mouse Initialize asynchronous buffers K_buf, M_buf WHILE recording is True: IF Key Event detected: Record timestamp t, event_type Push (t, event_type) to K_buf IF Mouse Event detected: Record timestamp t, X, Y coordinates Push (t, X, Y) to M_buf

**Algorithm 2: IKI/Digraph Feature Extraction** INPUT K_buf over window W Initialize IKI_list FOR i from 0 to length(K_buf) - 1: delta_t = K_buf[i+1].t - K_buf[i].t Append delta_t to IKI_list Calculate Mean(IKI_list), Variance(IKI_list) RETURN IKI_features

**Algorithm 3: CLC Computation** INPUT IKI_features, content_complexity_metrics Rank IKI_features -> R_IKI Rank content_complexity_metrics -> R_C Compute d_i = R_C_i - R_IKI_i for all i Compute rho_s = 1 - (6 * SUM(d_i^2)) / (n * (n^2 - 1)) RETURN rho_s as CLC

**Algorithm 4: Sigma-Lognormal Processing** INPUT M_buf Extract velocity profile |v(t)| from X, Y deltas Initialize lognormal primitives parameters (D, t_0, mu, sigma) Apply non-linear least squares optimization to fit primitives to |v(t)| RETURN optimized_primitives

**Algorithm 5: Signal Quality Validation** INPUT raw_velocity, reconstructed_velocity (from Alg 4) Calculate Residual Noise = raw_velocity - reconstructed_velocity Compute Signal Power / Noise Power Convert to Decibels (dB) IF SNR < 15: RETURN Status=Poor ELSE IF SNR < 25: RETURN Status=Good ELSE: RETURN Status=Excellent

**Algorithm 6: DSTER Inference** INPUT IKI_features Embed IKI_features into dense vectors Pass through Temporal-over-Channel Self-Attention block Pass through Channel-over-Temporal Self-Attention block Fuse stream outputs Compute Softmax probabilities RETURN Emotion_Label, Confidence_Score

**Algorithm 7: Bayesian Belief Update** INPUT prior_belief b_t, observation o_t, transition_matrix T, observation_matrix Z Initialize new_belief array FOR each state s_prime in S: sum_prob = 0 FOR each state s in S: sum_prob += T(s_prime | s) * b_t(s) new_belief[s_prime] = Z(o_t | s_prime) * sum_prob Normalize new_belief array so sum equals 1.0 RETURN new_belief

**Algorithm 8: Trust/Cognitive State Extraction** INPUT current_belief_state (15-dimensional array) Marginalize probabilities across the 5 Trust states Marginalize probabilities across the 3 Cognitive Load states RETURN dominant_trust_state, dominant_cog_state

**Algorithm 9: Adaptation Policy** INPUT dominant_trust_state, dominant_cog_state, task_risk IF dominant_cog_state == HIGH: Set verbosity = LOW Set UI_complexity = MINIMAL IF dominant_trust_state == LOW AND task_risk == HIGH: Set uncertainty_presentation = VISIBLE Determine target parameter shifts based on ruleset RETURN raw_adaptation_targets

**Algorithm 10: Dynamic Difficulty Adjustment** INPUT raw_adaptation_targets, current_parameters, smoothing_factor alpha FOR each parameter in current_parameters: new_param = (alpha * raw_adaptation_targets) + ((1 - alpha) * current_parameters) Clamp new_param to predefined [MIN, MAX] bounds RETURN bounded_parameters

**Algorithm 11: Hint Engine** INPUT dominant_cog_state, time_on_task Set threshold_1, threshold_2, threshold_3 based on dominant_cog_state IF time_on_task > threshold_3: RETURN "Explicit Solution" ELSE IF time_on_task > threshold_2: RETURN "Tactical Guide" ELSE IF time_on_task > threshold_1: RETURN "Subtle Nudge" ELSE: RETURN None

**Algorithm 12: HCXAI Explanation Generator** INPUT bounded_parameters, dominant_trust_state, DSTER_confidence IF DSTER_confidence < 0.6: explanation = "System is adapting based on general activity patterns." ELSE IF dominant_trust_state == HIGH: explanation = "Confidence in cognitive load is high; streamlining interface." ELSE: explanation = "Interface simplified to assist with current task complexity." RETURN explanation

**Algorithm 13: Complete Project ECHO Closed Loop** Initialize system variables, POMDP

matrices, and starting parameters WHILE Interaction Active: Execute Alg 1 (Capture) Execute Alg 2, 3, 4 (Extract Features) Execute Alg 5 (Validate SNR) Execute Alg 6 (DSTER Inference) Execute Alg 7 (Update Belief) Execute Alg 8 (Extract States) Execute Alg 9, 10, 11 (Determine Adaptation & Hints) Execute Alg 12 (Generate HCXAI) Push bounded_parameters and HCXAI to Application Interface Wait for next temporal window

# Complexity Analysis

Given the necessity for edge-deployment and minimal latency, computational complexity must be strictly bounded.

- **IKI & Correlation Computation:** Time complexity is \mathcal{O}(N \log N) due to the rank sorting required for the Spearman correlation over a temporal window of N keystrokes. Space complexity is \mathcal{O}(N).
- **Sigma-Lognormal Estimation:** The optimized extraction utilizes non-linear least squares fitting. Time complexity is approximately \mathcal{O}(M \times I) where M is the number of identified movement primitives and I is the number of optimization iterations.
- **Transformer Inference:** Standard transformer self-attention operates at \mathcal{O}(K^2 \cdot d), where K is the keystroke sequence length. Because ECHO utilizes short sliding windows, K remains sufficiently small, guaranteeing low inference latency.
- **Bayesian Belief Update:** Time complexity is \mathcal{O}(\vert{}S\vert{}^2) for matrix multiplication. Given the tightly constrained state space of \vert{}S\vert{} = 15, this resolves in \mathcal{O}(1) constant time for practical purposes, demanding a negligible CPU footprint.

# Test Case Design

The following illustrative test cases demonstrate expected system behavior under varying conditions.

| Test Case Scenario | Input | Expected Features | Expected Belief Shift | Expected Adaptation | Expected HCXAI Output *(Illustrative)* |
|---|---|---|---|---|---|
| **Normal typing** | Consistent cadence | IKI variance low, CLC = 0.40, SNR > 20 dB | Stable | None | *None* |
| **High hesitation** | Frequent long pauses | High IKI variance, high digraph delay | Slight shift to CogLoad: Medium | Pre-load hints | *None* |
| **High cognitive load** | Erratic typing, poor mouse sync | CLC > 0.60, SNR = 18 dB | Heavy shift to CogLoad: High | Decrease difficulty, simplify UI | "Task complexity appears high. Interface streamlined." |
| **Frustrated interaction** | Heavy keystrokes, rapid disjointed mouse | DSTER = Frustrated, Confidence > 0.8 | Shift to Trust: Low | Reduce enemy aggression | "Reducing encounter density." |

| Test Case Scenario | Input | Expected Features | Expected Belief Shift | Expected Adaptation | Expected HCXAI Output *(Illustrative)* |
|---|---|---|---|---|---|
| **Fatigue-like behavior** | Slowing velocity, trailing mouse curves | High Sigma-Lognormal variance (\sigma) | Shift to CogLoad: High, Trust: Med | Offer break, lower tension | "Performance tracking indicates fatigue. Suggesting a pause." |
| **Poor signal quality** | System lag, dropped inputs | SNR < 15 dB | None (Prior heavily weighted) | None | *None* |
| **High trust / high risk** | Consistent inputs during complex tasks | Trust = High, Risk flag active | Stable | Surface model uncertainty | "System confidence moderate; manual review recommended." |
| **Low trust / low risk** | Frequent manual overrides | Trust = Low | Shift to Trust: Very Low | Halt autonomous adjustments | "Automatic adjustments paused." |
| **Rapid success / mastery** | Fast, flawless inputs | DSTER = Engaged, CLC < 0.35 | Shift to CogLoad: Low | Increase difficulty gradually | "Scaling up challenge." |
| **Volatile user state** | Erratic switching between extremes | Confidence < 0.5 | Stable (Smoothing prevents oscillation) | None | *None* |

# Example Numerical Walkthrough

*Illustrative mathematical example — not a measured project result*. **Given Observation at t=1:**

- Mean IKI indicates high variance.
- CLC = 0.60 (High load indicator).
- SNR = 22 dB (Good signal).
- DSTER Output = Fatigued, Confidence = 0.75. **Bayesian Update Calculation:**
- Assume Prior b_0(s_{\text{high_load, low_trust}}) = 0.06.
- Transition probability T(s'\vert{}s) = 0.8 (assuming strong state persistence).
- Observation likelihood O(o\vert{}s') = 0.85 (fatigue and high CLC strongly align with high load).
- Unnormalized belief b'_{un} = 0.85 \times (0.8 \times 0.06) = 0.0408.
- Following normalization across the entire 15-state matrix, the posterior belief b_1(s_{\text{high_load, low_trust}}) rises significantly to 0.45. **Adaptation Calculation:**
- The policy engine determines the game parameter P_{target} should drop from 1.0 to 0.6.
- Applying exponential smoothing (\alpha=0.3): P_{new} = 0.3(0.6) + 0.7(1.0) = 0.18 + 0.70 = 0.88.
- Result: The game difficulty smoothly dips to 0.88 rather than jarringly crashing to 0.6, preserving immersion.

# Results

A. **Actual verified/project-derived results:** The core integration tests run on the simulator produced verifiable output: 59 passes, 3 failures, and 11 errors. These failures directly result from the known game_registry.py mapping mismatch outlined previously. B. **Synthetic validation results:** The DSTER model validates cleanly on artificially generated temporal clusters featuring strict mean and variance separations, validating the mathematical soundness of the codebase. C. **Proposed future evaluation:** The measurement of real-world DDA impact (e.g., NASA-TLX workload reductions, engagement metrics) relies on proposed expansive, multi-condition human-subject studies. *(Note: Real-world human efficacy is not demonstrated in the currently available project material.)*

# Project Contribution

The explicitly implemented contributions of Project ECHO include:

- A unified, reusable software-only HCI middleware architecture.
- The mathematical integration of Cognitive Load Correlation and Sigma-Lognormal parameters into a continuous, real-time pipeline.
- A fully realized 15-state Bayesian POMDP framework for continuous uncertainty tracking.
- A dynamic difficulty adaptation parameter modulation engine featuring bounded exponential smoothing.
- Comprehensive simulation and visualization utilities tailored for offline testing.

Proposed future contributions encompass the creation of vast, anonymized real-user interaction datasets, deployment-ready Unity and Unreal SDKs, and production-scale HCXAI frontends capable of seamless commercial integration.

# Final Research Thesis

Project ECHO operationalizes the following central thesis: *"Human-computer interaction systems can become more effective, empathetic, and reliable when interface adaptation is driven not solely by observable task performance, but by probabilistically estimated cognitive, affective, neuromotor, and trust-related user states derived non-intrusively."*
By integrating behavioral signal extraction with a Bayesian POMDP and a bounded HCXAI adaptation policy, ECHO effectively transforms static digital environments into cooperative, state-aware paradigms that respect and augment human autonomy.

# Conclusion

Project ECHO successfully establishes a comprehensive, software-only paradigm for human-centered adaptive computing. Moving decisively away from the physical intrusion of physiological biometric sensors, the architecture proves that mundane interaction telemetry—specifically keystrokes and mouse movements—contains rich latent representations of human cognition and affect. Through the rigorous implementation of the DSTER architecture, Sigma-Lognormal analysis, and a probabilistically sound Bayesian POMDP, ECHO ensures that system adaptations are not merely reactive, but contextually intelligent, stable, and bounded. As

the framework evolves toward commercial game engine integration and large-scale demographic validation across environments like Chennai's IT and educational sectors, it provides a foundational blueprint for empathic digital interfaces. Ultimately, Project ECHO adapts to the user's mental reality, safeguarding agency while optimizing the interactive experience.

## Works cited

1. DSTER: A Dual-Stream Transformer-based Emotion Recognition, https://www.researchgate.net/publication/385734166_DSTER_A_Dual-Stream_Transformer-based_Emotion_Recognition_Model_through_Keystrokes_Dynamics 2. (PDF) Neuromuscular Representation and Synthetic Generation of, https://www.researchgate.net/publication/267034507_Neuromuscular_Representation_and_Synthetic_Generation_of_Handwritten_Whiteboard_Notes 3. Sigma-Lognormal Modeling of Speech - PMC - NIH, https://pmc.ncbi.nlm.nih.gov/articles/PMC7943521/ 4. The Lognormality Principle and its Applications in e-Security, e, https://www.worldscientific.com/doi/pdf/10.1142/9789811226830_0001 5. Machine Learning and Deep Learning for Fixed-Text Keystroke, https://www.researchgate.net/publication/362029703_Machine_Learning_and_Deep_Learning_for_Fixed-Text_Keystroke_Dynamics